

Lauri Karttunen FestSchrift

Cleo Condoravdi (ed.)

April 22, 2018

**CENTER FOR THE STUDY
OF LANGUAGE
AND INFORMATION**

Contents

1	Integrating Implicative Commitments	1
	VALERIA DE PAIVA AND GERARD DE MELO	

Integrating Implicative Commitments

VALERIA DE PAIVA AND GERARD DE MELO

1.1 Introduction

Natural language processing systems, even when given proper syntactic and semantic interpretations, still lack the common sense inference capabilities required for genuinely understanding a sentence. Recently, there have been several studies developing a semantic classification of verbs and their sentential complements, aiming at determining which inferences people draw from them. Such constructions may give rise to implied commitments that the author normally cannot disavow without being incoherent or without contradicting herself, as described for instance in the work of Lauri Karttunen. In this paper, we model such knowledge at the semantic level by attempting to associate such inferences with specific word senses, drawing on WordNet and VerbNet. This allows us to investigate to what extent the inferences apply to semantically equivalent words within and across languages.

Understanding what a text means requires more than just decoding its syntactic and semantic structure. Rather, it typically involves drawing additional inferences beyond the information that it immediately conveys. Even when supplied with proper syntactic and semantic interpretations of a sentence, current natural language processing systems still lack the common sense inference capabilities required to interpret it in the way humans do. Given a sentence such as “John missed that Mary had left”, we are inclined to infer that Mary had indeed left. In contrast, given the sentence “John pretended that Mary had left”, we

are inclined to suppose the opposite. Although both sentences share a common structure, involving verbs with sentential complements, there are clear differences in the types of implicative commitments the author is making. Such sentences have been studied in detail in the work of Karttunen.

We have been working on inferential properties of linguistic expressions to enable automated reasoning for natural language understanding. In particular, we are interested in linguistic expressions that can be said to ‘flip’ (e.g. change from true to false) or ‘erase’ (if X is a true proposition, ‘Mary said X ’ can be true or false) the truth values naturally associated with some other expressions. For instance, hearing that “John prevented Mary from leaving”, we are inclined to suppose that Mary has not left and that there might be reasons for this not leaving.

Textual inferences may be based on purely linguistic knowledge, assumptions about language use by collaborative rational agents, knowledge about the world, or any combination thereof, and it is important to recognize how much of this information can be codified in the lexicon.

The semantics of complement constructions is an important part of textual inference. Karttunen’s seminal work on English complement constructions (Karttunen, 1971) makes the point that while it is valid and helpful to classify verbal constructions according to their syntactic characteristics, e.g., whether they take propositional complements, it is also valid and even more useful to classify them according to their semantic characteristics, such as the *factivity* or *implicativity* of complements. Karttunen’s more recent work (with Condoravdi and Nairn (Nairn et al., 2006), and, independently with Zaenen (Zaenen and Karttunen, 2013)) takes a more empirical route and discovers the factivity and implicativity behavior by examining corpora on the web. We follow this tradition, aiming for as broad coverage as possible. As remarked in Nairn et al. (2006) this type of semantic information is not available in, or deducible from any public lexical databases such as WordNet (Fellbaum, 1998), VerbNet or FrameNet, by itself. Hence, they had to compile their own tables of “implication signatures” for a large class of complement-taking constructions. To create these tables, Nairn, Condoravdi and Karttunen write

“[we took] the verbs in order of decreasing frequency in the British National Corpus (BNC), we determined their natural implications (if any). Judgments were based on agreement by multiple annotators using resources such as Google search and the Linguist’s Search Engine to sample the relevant constructions in the wild.”

Of the 1,250 relevant verbs in that sample, they classified 400 in a first pass and ended up with 108 involving `that` complements (`compex-that`) and 114 with infinitive complements (`xcomp-inf`). The original lists are available from CSLI at http://web.stanford.edu/group/csl_i_lnr/Lexical_Resources/.

Our hypothesis is that implicative commitments derive from the semantics of the involved words and hence are somewhat independent of the particular choice of word among a set of synonyms, certainly grammatical theory independent, and perhaps to some degree also language independent. While there are certainly exceptions to this, our system can only get better as we provide more relevant data points to it so as to enable more accurate generalizations.

Based on this hypothesis, our previous work (de Melo and de Paiva, 2014) took the marked implicative and factive verbs from Karttunen et al., as distributed by CSLI, and made them more semantic, by considering the various meanings of these verbs and marking in the Princeton WordNet (Fellbaum, 1998) those synsets that corresponded to the relevant senses of the verbs.

The work in this paper extends the original lexical resource into a system that is able to generalize beyond the specific markings found in the lexicon. Our system operates on universal dependencies (Nivre et al., 2016) and uses semantic relatedness assessments to infer information for words not in the lexicon. To make this system more complete, we have also integrated information from additional sources beyond what was in our initial resource. These include implicative commitments based on adjectives and context-creating nouns. We have also worked with annotators to obtain initial lists of markings for Brazilian Portuguese.

1.2 Related Work

Such differences have been examined in detail by recent studies that have attempted to develop a classification of verbs (and verb-noun collections) that take sentential complements (Karttunen, 2012). Similar analyses can be made with respect to other constructions, e.g. to study adjective constructions like “It is confusing that Mary has left” vs. “It is improbable that Mary has left” (Zaenen and Karttunen, 2013). Likewise, one can also study implicative commitments regarding the existence of entities. For instance, “She cancelled the meeting” generally leads us to believe that there was no meeting, while “They caused a strike” means that a strike did occur (Amaral et al., 2012).

Certain sentences give rise to implied commitments that the author

normally cannot disavow without being incoherent or without contradicting herself. Previous work has developed detailed classifications of words and their implicative commitments at the lexical level (Nairn et al., 2006). In this paper, we build on this work but attempt to model such knowledge about commitments at the semantic level instead. For this, the commitments are tied to specific word senses, drawing on WordNet and VerbNet as sense inventories. We rely on automatic disambiguation methods to produce a lexicon of sense-specific implicative commitments. Among other things, this allows us to investigate to what extent such inferences also apply to semantically equivalent words within and across languages. Our findings indicate that they are for the most part preserved when transitioning from words to synonym sets, and at least most of the English classifications collected by Nairn et al. (Nairn et al., 2006) seem to transfer to other languages such as Portuguese as well.

1.3 Preliminaries

Karttunen’s seminal “The Logic of English Predicate Complement Constructions” (Karttunen, 1971) makes the point that while it is valid and helpful to classify verbal constructions according to their syntactic characteristics, e.g. whether they take propositional complements, it is also valid and even more useful to classify them according to their semantic characteristics, such as the *factivity* of complements. The term *factive verb* introduced by Kiparsky and Kiparsky (Kiparsky and Kiparsky, 1970) refers to the notion that any simple assertion using one of these verbs (e.g. “John knew that Mary had left”) commits the speaker to the belief that the complement sentence (“Mary had left”), just by itself, is also true.

Karttunen observes that it is sometimes “possible to show that there is a definite connection between the semantic properties of a verb and certain syntactic characteristics”, and thus goes on to investigate parts of this connection. The syntactic characteristics referred to are that these are propositional complement constructions, while the semantic characteristics are what he refers to as the “logic” in the title of his book. The latter is more general than simply factivity. In fact, 35 years later, as part of a project to extract information from text for question answering, Karttunen, together with Nairn and Condoravdi (Nairn et al., 2006) presented an algorithm for detecting author commitment to the truth or falsity of complement clauses based on their syntactic type and on the meaning of their embedding predicate. They also created a small lexicon of around 300 verbs spelling out the implicative

commitments that these verbs indicate. This was the starting point of our work. We contend that such implicative commitments are best described at the semantic level, as one would hope given Karttunen’s choice of the word “logic” in the title of his seminal work.

The same applies to other sources and forms of implicative commitments. For instance, in some cases, nouns can induce factivity commitments. As an example, “John took the opportunity to sing” entails that he did sing (Price et al., 2008). Similarly, adjectives, too, can be marked for factivity and implicativity, as in a recent study (Zaenen and Karttunen, 2013) that collected such annotations using Amazon Mechanical Turk (<https://www.mturk.com/>). Additionally, one can study implicative commitments relating to existential and temporal aspects (Amaral et al., 2012).

Throughout this paper, we are adopting a generalized, somewhat underspecified notion of implicativity that subsumes factivity, entailment, and presupposition, calling them “implicative commitments” for want of a better name. This notion as described by Nairn et al. (Nairn et al., 2006) seems more robust and clear-cut, as exemplified in their small lexicon. So no attempt will be made in the following to distinguish between the several components of the “implicative commitments”.

However, it is clear that our lexicon would have to be just one of several ingredients in any natural language understanding systems aimed at fully assessing the *factuality* of textual content, as defined by Saurí and Pustejovsky’s in Saurí and Pustejovsky (2009, 2008, 2012). We might want to know whether an event mentioned in the text has really occurred in the world. Who is the source of the information? What is the stance of the author of the text? Does the author indicate whether he or she believes the source?

This encompassing problem is also discussed by Karttunen and Zaenen in Karttunen and Zaenen (2005). They call “veridicity” what we can ascertain from the text itself about the stance of the author as far as her text is concerned. As they explain

When somebody says or writes *The earth is flat* or *The king of France is bald* because (s)he is a liar or ill-informed, nothing in these linguistic expressions in themselves alerts us to the fact that they do not correspond to situations in the real world. But a text does give us information about the stance its author takes vis-à-vis the events or states described. When somebody says *Bill acknowledges that the earth is round*, we know something about the author’s as well as Bill’s beliefs in the matter.

A different perspective is taken in Saurí and Pustejovsky’s FactBank corpus and factuality studies (Saurí and Pustejovsky, 2009, 2008, 2012).

The work in FactBank is an attempt to mark the implicative commitments of **both the author and some participants of the events**, in some given text. As Saurí and Pustejovsky explain

Identifying the veracity, or factuality, of event mentions in text is fundamental for reasoning about eventualities in discourse. Inferences derived from events judged as not having happened, or as being only possible, are different from those derived from events evaluated as factual. Event factuality involves two separate levels of information. On the one hand, it deals with polarity, which distinguishes between positive and negative instantiations of events. On the other, it has to do with degrees of certainty (e.g., possible, probable), an information level generally subsumed under the category of epistemic modality.

Thus their work is not concerned with what the author’s stance with respect to his/her text, but with some over-arching version of the actuality of events, as determined by the reader.

The work on FactBank has been extended by the Stanford NLP group (de Marneffe et al., 2011), who re-annotated the same sentences using workers on Amazon Mechanical Turk, given their stance that textual entailment is to be measured in terms of the understanding of text by ordinary people, not the one of logicians, philosophers, or linguists. De Marneffe, Manning, and Potts explain (de Marneffe et al., 2011)

Natural language understanding depends heavily on assessing veridicality — whether the speaker intends to convey that events mentioned are actual, non-actual, or uncertain. However, this property is little used in relation and event extraction systems, and the work that has been done has generally assumed that it can be captured by lexical semantic properties. Here, we show that context and world knowledge play a significant role in shaping veridicality. We extend the FactBank corpus, which contains semantically driven veridicality annotations, with pragmatically informed ones. Our annotations are more complex than the lexical assumption predicts but systematic enough to be included in computational work on textual understanding. They also indicate that veridicality judgments are not always categorical, and should therefore be modeled as distributions.

Thus following Karttunen and Zaenen, we make a distinction between factuality (what has actually happened in the world) contrasting it with veridicity (what the author of the utterance conveys about what has happened in the world). While we agree that in the long run we will need distributions of veridicality judgements for sentences that require opinions and assessing the trustworthiness of sources, we believe that there is a level of assessment that can be coded up in the lexicon. It is

rarely disputed that if a text says “John admitted that he was guilty” then the author of the text believes that the embedded clause “John is guilty” is true. Whether the reader of the text should or should not believe the text he is reading (the veridicality of the text) is an orthogonal question. We are only concerned with the fact that by using the predicate “admit” the author of the sentence indicates that his stance is that John was guilty. That is, we are concerned with the veridicity of the sentence, not its veridicality, which is much harder to ascertain. In particular veridicality of sentences depends on the veridicity of authors that are not lying and that are well-informed. Deciding whether authors are not lying and whether they are well-informed (or not) is much harder than deciding what the linguistic text tells us.

While veridicality or factuality is a possible long term goal, we instead start from the less lofty goal of *veridicity*, a level of assessment that we believe can be coded up in the lexicon. As Karttunen and Zaenen put it:

It is useful to distinguish between two ingredients that go into determining the truth value of utterance, one is the trustworthiness of the utterer and the other is the stance of the utterer vis-à-vis the truth of the content. The latter we will call the veridicity of the content [...].

The classification chosen by Nairn et al. (Nairn et al., 2006) focuses on cases in which the author’s commitment to the truth of a complement clause arises solely from the larger sentence it belongs to, leaving aside other sources of information about the beliefs of the author. Note that it can be difficult to decide between *entailments*, that is, what the author is actually committed to, and the more pragmatic notion of *conversational implicatures*, that is, what a reader/hearer may feel entitled to infer. For example, “Ed did not refuse to participate” might lead the hearer to conclude that Ed participated. But the speaker could continue with “He was not even eligible” indicating the opposite.

The tags used to describe the implicative commitments of lexical items need to distinguish several types of context-specific behaviors. Some words yield an entailment in both affirmative and negative environments but there are others, “one-way implicatives”, that yield entailments only in one or the other environment. Furthermore, the entailment may be either positive or negative depending on the polarity of the environment. Altogether Nairn et al. provide us with a table of “implication signatures” for a large class of complement-taking constructions. These implication signatures can be stacked together, which gives rise to their algorithmic solution for dealing with more complex phrases (Nairn et al., 2006).

1.4 Integrating Annotations

In this paper we advance the state-of-the-art in that, instead of associating implicative commitments just with lexical items, we associate them with word senses, as represented by synsets given by the lexical resource WordNet (Fellbaum, 1998). These subsequently, in Section 1.5, shall allow us to relate new words to the word senses covered in the lexicon.

Our input is a set of resources that describe implicative commitments at the lexical level (included the Nairn et al. lexicon and further resources listed later on in Section 1.6.1). Our output will contain descriptions of implicative commitments at the sense level.

Generally speaking, an input entry at the lexical level consists of a partially instantiated construction and a set of tags describing the implicative commitments of that construction in particular contexts. The construction may involve one or more lexical items to be disambiguated. For instance, often, there is simply a single verb like “to find” in a syntactic frame expecting a subject and a that-complement (“John found that Mary had arrived”). Instead of describing the construction for a specific lexical item like “to find”, we can instead characterize it at a more abstract semantic level, allowing an instantiation with several additional words (“notice”, “discover”, etc.). We thus have to disambiguate the word “to find” in order to establish which senses are relevant and know which sense-specific synonyms can be assumed to give rise to the same commitments.

In some cases, the construction may also involve two lexical items that need to be disambiguated. For instance, the construction may involve both the verb “avoid” and a description of its possible arguments, described by the word “event”. These will then both be mapped to WordNet synsets, and WordNet’s synset hierarchy will allow us to recognize specific instantiations, e.g. “avoiding a strike” vs. “avoiding a road bump” – In the first case, there is no strike, whereas in the second case, the road bump remains.

1.4.1 Sense Mappings and Disambiguation

Our goal is thus to map one or more input lexical items $e = (l, p, f)$ on the one side, each consisting of a lemma l , part-of-speech tag p , and syntactic frame marker f , to one or more candidate WordNet synsets s on the other side. For example, our input entry could be for the lemma “admit” with part-of-speech tag **verb**, and a marker for a syntactic frame with infinitival complement (as in “The report was admitted to be incorrect”). WordNet lists a total of 8 verb senses for the input

lemma “admit”. For instance, the first WordNet sense for “admit” refers to the process of allowing people to enter a building, and the second one refers to admittance into a group or a community. These do not reflect the meaning that our input entry is targeting. Our system instead needs to choose the WordNet sense for *declaring to be true or admitting the existence or reality or truth of something*.

In order to achieve such disambiguated mappings in the case of verbs, we draw on another traditional lexical resource VerbNet (Kipper-Schuler, 2005). VerbNet is a lexicon of Levin-style verb classes (Levin, 1991) based on the syntactic frames of English verbs. Since the input lexicons often refer to specific syntactic frames, e.g. for sentential complements, the syntactic characterizations present in VerbNet can help us pick relevant senses.

For a given input entry e with lemma l , let $W(l)$ denote the set of WordNet senses and $V(l)$ denote the set of VerbNet senses for l . We first use the following similarity function to score the candidate VerbNet senses $v \in V(l)$:

$$\text{sim}(e, v) = \frac{1}{|V(l)|} \left(\epsilon + \max_{f_v \in F(v)} \text{sim}(f_v, f) \right) \quad (1.1)$$

Here, the strength of association between e and one of the verb senses $v \in V(l)$ first of all depends on the number of candidates $|V(l)|$. Additionally, we obtain the relevant VerbNet class and its syntactic frame information from VerbNet, using $F(v)$ to denote the set of syntactic frames associated with the VerbNet class of v .

We rely on a simple rule-based similarity measure between syntactic descriptors f_v from VerbNet and descriptors f from the original lexicon, which in our case are based on PARC’s Bridge system (Bobrow et al., 2007). For instance, VerbNet’s NP **V that S** would match the **V-SUBJ-COMPLEXthat** descriptor from the Bridge system but not **V-SUBJ-OBJexpl-XCOMPinf**. While this similarity function clearly prefers VerbNet senses that have at least one matching syntactic frame, we still retain the remaining VerbNet senses as possible candidates by setting ϵ to a small constant greater than 0, because VerbNet’s syntactic frame descriptions are not always complete.

We then use these VerbNet similarities to help us in scoring the candidate WordNet synsets $s \in W(l)$:

$$\text{sim}(e, s) = \begin{cases} \epsilon + \sum_{v \in V(l)} \frac{\max_{s_v \in V_w(v)} \text{sim}(s_v, s)}{1 + |V_w(v)|} \text{sim}(e, v) & s \in W_p \\ 0 & s \notin W_p \end{cases} \quad (1.2)$$

Here, first of all, for a given part-of-speech tag p , W_p denotes the global set of all WordNet senses for that part-of-speech tag. If we are attempting to disambiguate an entry e for which the part-of-speech tag p denotes adjectives, for instance, then only adjective synsets $s \in W_p$ receive a non-zero similarity score.

For a given VerbNet sense v , we use $V_w(v)$ to denote the set of WordNet synsets mapped to v . For this, we make use of the existing disambiguated mappings between VerbNet and WordNet provided with the former. VerbNet’s verb senses tend to be more coarse-grained, as their induction was mainly guided by syntactic considerations. Thus for a given VerbNet sense v , the set $V_w(v)$ may contain multiple WordNet synsets. We compare these synsets with our current synset of interest s , simply by using the identity function as the similarity measure $\text{sim}(s_v, s)$, although there are also many existing WordNet sense similarity measures that could be plugged in here. If none of the synsets $s_v \in V_w(v)$ are similar, then the overall value of $\text{sim}(e, s)$ will turn out to be just a small $\epsilon > 0$. In contrast, if there are similar synsets, we discount using the set cardinality $|V_w(v)|$, based on the intuition that the score should be higher if there are fewer possible synset matches and thus less ambiguity.

These scoring functions thus allow us to select likely WordNet synsets for a given input entry.

1.4.2 Lexicon

Once we have disambiguated all the lexical items involved in the construction based on the disambiguation procedure described above, we can assign the implicative commitment tags associated with the original construction to our newly sense-disambiguated version of the construction. The latter refers to specific senses of lexical items, thus eliminating possible confusion about which senses are meant in the case of ambiguous words. At the same time, the latter is more general by enabling us to instantiate the construction with alternative synonymous words, rather than forcing us to select the original lexical item.

In addition to the new sense-specific implicative commitment markings, our final lexicon incorporates information from WordNet, VerbNet, and other lexical resources. For a given word, we list a number of (potentially overlapping) senses. Each sense corresponds either to a WordNet synset or to VerbNet sense or to both, based on the mappings between the two resources provided with VerbNet.

For some of these word senses, we now have information about constructions they can be involved in and the implicative commitments that are entailed when relevant words with those senses (or hyponym

senses) are used in those specific constructions. The implicative commitments are provided as properties of the constructions, marked using tags like **fact.p** (based on the ones in the original lexicons). The constructions are linked to specific word senses and thus our lexicon describes the implicative commitments that words with a corresponding word sense entail when used in the corresponding syntactic constellation.

Additionally, our lexicon also provides references to other resources. For every sense corresponding to a VerbNet entry, we list the corresponding FrameNet (Baker et al., 1998) frames, relying on the Sem-Link project’s mappings between VerbNet and FrameNet (Bonial et al., 2013). For every sense corresponding to a WordNet entry, we also include the relevant SUMO (Niles and Pease, 2001) concepts based on the pre-existing WordNet-SUMO mappings (Niles and Pease, 2003).

We make use of the RDF standard for information interchange, but our data can be converted to any number of other formats, including simple tab-separated-value files. It is clear that this is just a first step towards a freely available open “unified” lexicon in the spirit of Crouch and King (Crouch and King, 2005). In future work, we additionally hope to add presupposition relationships between different verbs, as studied by Temper and Frank (Temper and Frank, 2013) (e.g., given “Spain *won* the tournament”, one may presuppose that Spain *played* in the tournament). Another possible extension is to mine corpora for information about downward entailment as proposed by Danescu-Niculescu-Mizil et al. (Danescu-Niculescu-Mizil et al., 2009).

1.5 Implicative Commitments for Sentences

Our current work extends the original lexical resource into a natural language understanding system that is able to generalize beyond the specific markings found in the lexicon.

1.5.1 Syntactic Analysis

Our system starts from syntactic analyses in terms of Universal Dependency annotations (<http://universaldependencies.org/>).

Our system needs to recognize what in more theoretical work we called *context markers*. Some verbs that denote context markers are not as informative as factives or implicatives. Instead of assigning a new truth value to the propositions that are nested into the contexts they introduce, they simply make these contexts opaque. Thus if I say *John opened the door*, you may conclude that I think that John opened the door, assuming that I do not usually lie. If I say *John managed to open the door* you may wonder whether there was a problem that

made it difficult to open the door or not, but you can still conclude that *John opened the door*. And if I say *Mary prevented John from opening the door* you may rationally conclude that *John opened the door* is false, i.e John did not open the door. However, for a sentence like *Mary said that John opened the door*, the stance of the author remains unclear. We do not know if they believe Mary to be a well-informed and honest person and hence believe that John did open the door, or if they do not believe the stated fact, but do not want to say it directly. There are many synonyms for the verb *to say*, and these introduce a context in the representation that one could decide to consider factive, for instance, if one considers a source such as *The New York Times* as always reputable and reliable. In order to recognize syntactic context-introducing verbs that do not allow for any clear inference, e.g., for the verb *to say*, we extend the lexicon with relevant markings of such verbs. Subcategorization frame information helps in this process, but is under-realized both in Princeton WordNet and in VerbNet. Thus, we extended the lexicon by marking around a thousand instances in our Implicative Lexicon.

1.5.2 Prediction

Given an identified context-inducing word, we rely on our lexicon to infer relevant implicative commitments. For words covered by the lexicon, it is just a matter of recognizing them within the syntactic structure of the sentence. When a new word, not covered by the lexicon, is encountered, we rely on a measure of semantic similarity to find relevant points of reference in the lexicon that allow us to guess the implicative behavior of the new word.

To this end, we rely on the Open Multilingual WordNet (Bond and Foster, 2013), which combines data from a number of multilingual WordNet-like projects, including the Princeton WordNet for English (Fellbaum, 1998) and our own OpenWN-PT (Rademaker et al., 2014). For a given target word, we retrieve its list of word senses from the multilingual WordNet data. Then, we rely on the Wu-Palmer semantic relatedness measure between synsets to relate these senses to the entries in our lexicon. The output markings are a weighted sum of the markings of the nearest neighbours.

In the long run, we endeavor to build on our previous work to lay the basis of a multilingual resource and system that can help us to achieve better and more fine-grained inferential results than the ones we have for recognizing textual entailment in English nowadays. As the accuracy of most textual entailment recognition systems remains below 60% (Zanzotto and Pennacchiotti, 2010), there is much room for

improvement.

1.6 Results

1.6.1 Data

For our input, we relied on the lexicons from Nairn et al. (Nairn et al., 2006) and related papers, as distributed by CSLI¹, the creation of which was largely DARPA-funded, in the context of the AQUAINT project.

A lexicon entry consists of a lemma, a syntactic category descriptor, one or more tags to describe the implicative commitment behavior, and frequently also an example sentence for the lemma in the respective syntactic frame highlighting such a possible implicative commitment.

In addition to the simple factives data and the simple implicatives data, we also included the more recent phrasal implicatives data (Karttunen, 2012), which captures implicative behaviors that depend on the specific arguments of a verb. We likewise included the CSLI data about the factivity of adjectival statements, e.g. “It is accurate that John informed the president” vs. “It is untrue that John informed the president, as well as the reduced set of markings on nouns from the appendix of Price et al. (Price et al., 2008).

Moreover, we incorporated the CSLI infinite temporal markings data, which provides information about possible temporal shifts in addition to factivity information. For instance, for a sentence like “Mary persuaded Ed to cook dinner”, the cooking can happen after the persuasion, while for “Mary let Ed cook dinner” the cooking typically is simultaneous with the letting. Finally, we created a lexicon of entries about existential commitments, drawing on analyses in the Amaral et al. paper (Amaral et al., 2012), among other sources.

To disambiguate all of this data with respect to WordNet, we made use of the algorithm described in Section 1.4.1, choosing the highest-ranked synset for each lexical item in the input. In the case of ties (in particular for the nouns and adjectives, where the VerbNet-based heuristics do not apply), we chose the top WordNet synset in terms of WordNet’s original frequency-based sense ranking.

Table 1 summarizes the results of this process. For each category of information, we describe the number of lexical items in the input and the relevant number of candidate senses in VerbNet and WordNet. We then list the number of annotated entries in the input and in our disambiguated output. Since for many data types, annotated entries only refer to a single lexical item, these numbers are often quite correlated. On a random sample of 20 simple implicatives, our automatic disam-

¹http://www.stanford.edu/group/csl_i_lnr/Lexical_Resources/

TABLE 1 Implicative Commitment Data

Data Type	Lexical Items			Annotated Entries	
	Input	VN		Input	Disamb.
		Cand.	WN		
Simple factives	108	188	486	108	108
Simple implicatives	114	223	950	114	114
Phrasal implicatives	17	67	584	14	14
Factiv. of adj. comp.	278	0	449	278	225
Factiv. of extraposed adj. comp.	695	0	1,534	695	623
Factiv. of noun comp.	112	0	304	112	108
Temporal implicatives	86	160	811	86	86
Existential commitments	74	101	428	62	62

biguation had a precision of 75.0%. This is a reasonable result given that we are using fine-grained WordNet word senses and that particularly for verbs we have to choose from a very large set of candidates, as can be seen in Table 1. We could not reliably assess the recall of our disambiguation, because for many of the rarer word senses, it is more difficult to establish whether the same implicatives hold as for the primary senses of those lexical items.

In any case, we are currently in the process of relying on human annotators to correct and extend our data, giving us a more complete set of reliable WordNet sense annotations.

1.6.2 Discussion

One of the low hanging fruits of using our sense-specific implicative markings is the extension of coverage it affords. For example the original lexical markings repository provides the marking for “to bother” but not for its WordNet synonym “to trouble”, as in “It troubled Ed that Mary was about to leave”. Similarly there is a marking for “to acknowledge” as in “They acknowledged that the report was correct” but not for “to concede” in the same frame.

Obviously, not all synonyms can be used in the syntactic frame specified by the annotation entry. WordNet’s synonym sets are semantically motivated, and as such may include words with different syntactic behavior. Thus the extension only applies to words with compatible syntactic behavior. For example, “to hate” is marked as factive as in “Ed hated to leave the party” or “Ed hated that Mary went home”. The synonymous verb “to detest” was not marked in the original lexicon, but is a possible synonym listed by our lexicon.

Another advantage of a semantic approach is that we can more ac-

curately describe the behavior of words when this behavior depends on the semantic context. For example, if you have “wasted your chance to go to Paris”, then you probably did not go, whereas if you “wasted your money to go to Paris”, then you probably did go (Karttunen, 2012). Similarly, “He withdrew his hand” does not affect the existence of the hand, while “He withdrew his offer” does reflect a change regarding the existence of the offer. Our lexicon distinguishes the two cases based on the semantic type of the complement. Given a novel example, an NLP system can make use of WordNet’s hypernym hierarchy or of WordNet-based semantic relatedness measures to determine which of the two cases is more likely.

Of course, our lexicon nevertheless lacks descriptions to account for certain more involved contexts. It seems difficult to account for all possible context-specific behaviors. For example, while “to cause something” generally implies that that something comes to be, there are also examples like “The decree was causing a revolution when it was revoked”, from which one is likely to conclude that a revolution was ultimately avoided (Amaral et al., 2012).

1.6.3 Cross-Lingual Applicability

Our WordNet sense markings not only enable us to find new synonyms, but also allow us to look up non-English equivalents in sense-aligned non-English versions of WordNet, such as EuroWordNet (Vossen, 1998) and UWN (de Melo and Weikum, 2009).

We have additionally arranged for a human-created Portuguese translation of the Nairn et al. lexicon and examples², and checked that most of the inferential behavior is preserved under direct translations. Some anecdotal observations can be made about this translation. First, as expected, the translation tends to indicate more English verbs, with subtle variations of meaning, going to a single Portuguese verb, for example “abhor” and “abominate” mapping to “abominar”, or “acknowledge” and “recognize” to “reconhecer”, or, more telling, “amaze”/“astonish”/“surprise” going to “surpreender”. All the factive verbs and their examples provided by Nairn et al. seem to work in Portuguese, very much like in English. This is despite some non-direct translations, for example sometimes a single work in English like “affect” becomes a phrasal verb in Portuguese “fazer de conta” (make believe), or “to perplex” becomes “deixar perplexo”. Things are not always as clear-cut on the implicatives and their examples. More work is needed here.

²The authors would like to thank Henrique Oliveira for his help with the Portuguese translations.

1.7 Conclusion

In this paper, we modeled the common sense inference capabilities required for genuinely understanding a sentence as knowledge at the semantic level by associating such inferences with specific word senses, drawing on WordNet and VerbNet. This allowed us to determine to what extent the inferences apply to semantically equivalent words within and across languages. Our preliminary lexical resources are available from <http://lexvo.org/implicative-lexicon/>.

Much work remains to be done, both for multiple languages and for extending the kinds of inferences we can deal with. We are actively integrating additional data. Incorporating new data resources requires aligning and integrating their entries into the Implicative Lexicon. We are also drawing on the large-scale entailment and contradiction detection (ECD) corpora that have appeared in recent years (Marelli et al., 2014, Bowman et al., 2015). Other future sources include the work of Lin and Pantel (Chklovski and Pantel, 2004), of Burchardt and Frank (Burchardt and Frank, 2006) of Tremper and Frank (Tremper and Frank, 2011, 2013) and of Dagan and collaborators (Lotan et al., 2013, Stanovsky et al., 2017). We continue to expand the Implicative Lexicon to include implicative commitments based on adjectives and context-creating nouns (Gurevich et al., 2008). The implicative commitments of adjectives and nouns are less understood than those of verbs and we hypothesize that our knowledge-based approach will help to abstract across different parts of speech.

Finally, by making our lexicon and techniques available to the research community, we hope to encourage research into and the creation of similar resources for multiple languages. Our efforts to develop the Implicative Lexicon for Brazilian Portuguese indicate that English resources can be used to bootstrap Brazilian Portuguese ones but cannot be directly ported from one language to another.

References

- Amaral, Patricia, Valeria de Paiva, Cleo Condoravdi, and Annie Zaenen. 2012. Where's the meeting that was cancelled? existential implications of transitive verbs. In *Proceedings of the 3rd Workshop on Cognitive Aspects of the Lexicon*, pages 183–194. Mumbai, India: The COLING 2012 Organizing Committee.
- Baker, Collin F., Charles J. Fillmore, and John B. Lowe. 1998. The Berkeley FrameNet project. In *Proc. COLING-ACL 1998*, pages 86–90.
- Bobrow, Daniel G., Bob Cheslow, Cleo Condoravdi, Lauri Karttunen, Tracy Holloway King, Rowan Nairn, Valeria de Paiva, Charlotte Price, and Annie Zaenen. 2007. PARC's Bridge and Question Answering sys-

- tem. In T. H. King and E. M. Bender, eds., *Proceedings of the GEAF07 Workshop*, pages 46–66. Stanford, CA: CSLI.
- Bond, Francis and Ryan Foster. 2013. Linking and extending an open multilingual wordnet. In *ACL (1)*, pages 1352–1362.
- Bonial, Claire, Kevin Stowe, and Martha Palmer. 2013. Renewing and revising SemLink. In *Proc. 2nd Workshop on Linked Data in Linguistics*. Pisa, Italy.
- Bowman, Samuel R., Gabor Angeli, Christopher Potts, and Christopher D. Manning. 2015. A large annotated corpus for learning natural language inference. In *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Association for Computational Linguistics.
- Burchardt, Aljoscha and Anette Frank. 2006. Approaching textual entailment with lfg and framenet frames. *Proc. of the Second PASCAL RTE Challenge Workshop*.
- Chklovski, Timothy and Patrick Pantel. 2004. Verbocean: Mining the web for fine-grained semantic verb relations. In *Proceedings of EMNLP, vol 4*.
- Crouch, D. and T.H. King. 2005. Unifying lexical resources. In *Proceedings of the Interdisciplinary Workshop on the Identification and Representation of Verb Features and Verb Classes*. Saarbruecken, Germany.
- Danescu-Niculescu-Mizil, Cristian, Lillian Lee, and Richard Duccott. 2009. Without a “doubt”? Unsupervised discovery of downward-entailing operators. In *HLT-NAACL*, pages 137–145. The Association for Computational Linguistics. ISBN 978-1-932432-41-1.
- de Marneffe, Marie-Catherine, Christopher D. Manning, and Christopher Potts. 2011. Veridicality and utterance meaning. In *Proceedings of the Fifth IEEE International Conference on Semantic Computing: Workshop on Semantic Annotation for Computational Linguistic Resources, Stanford, CA*. IEEE Computer Society Press.
- de Melo, Gerard and Valeria de Paiva. 2014. Sense-specific implicative commitments. In *Computational Linguistics and Intelligent Text Processing*, pages 391–402. Springer Berlin Heidelberg.
- de Melo, Gerard and Gerhard Weikum. 2009. Towards a Universal Wordnet by learning from combined evidence. In *Proceedings of the 18th ACM Conference on Information and Knowledge Management, CIKM '09*, pages 513–522. New York, NY, USA: ACM. ISBN 978-1-60558-512-3.
- Fellbaum, Christiane, ed. 1998. *WordNet: An Electronic Lexical Database*. The MIT Press. ISBN 026206197X.
- Gurevich, Olga, Richard Crouch, Tracy Holloway King, and Valeria de Paiva. 2008. Deverbal nouns in knowledge representation. *Journal of Logic and Computation* 18:385–404.
- Karttunen, Lauri. 1971. *The logic of English predicate complement constructions*. Bloomington, IN: Publications of the Indiana University Linguistics Club.

- Karttunen, Lauri. 2012. Simple and phrasal implicatives. In *Proceedings of the First Joint Conference on Lexical and Computational Semantics - Volume 1: Proceedings of the Main Conference and the Shared Task, and Volume 2: Proceedings of the Sixth International Workshop on Semantic Evaluation*, SemEval '12, pages 124–131. Stroudsburg, PA, USA: Association for Computational Linguistics.
- Karttunen, Lauri and Annie Zaenen. 2005. Veridicity. In *Annotating, Extracting and Reasoning about Time and Events. Dagstuhl Seminar Proceedings 05151*, G. Katz, J. Pustejovsky, F. Schilder (eds.), Dagstuhl. Germany: Dagstuhl, Germany.
- Kiparsky, Paul and Carol Kiparsky. 1970. Fact. In M. Bierwisch and K. E. Heidolph, eds., *Progress in Linguistics: A Collection of Papers*, pages 143–73. The Hague: Mouton.
- Kipper-Schuler, Karin. 2005. *VerbNet: A Broad-Coverage, Comprehensive Verb Lexicon*. Ph.D. thesis, University of Pennsylvania.
- Levin, Beth. 1991. *English Verb Classes and Alternations: A Preliminary Investigation*. University of Chicago Press.
- Lotan, Amnon, Asher Stern, and Ido Dagan. 2013. Truth-teller: Annotating predicate truth. In *Proceedings of the 2013 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 752–757.
- Marrelli, Marco, Stefano Menini, Marco Baroni, Luisa Bentivogli, Raffaella Bernardi, and Roberto Zamparelli. 2014. A SICK cure for the evaluation of compositional distributional semantic models. In *Proceedings of LREC 2014*.
- Nairn, Rowan, Cleo Condoravdi, and Lauri Karttunen. 2006. Computing relative polarity for textual inference. *Inference in Computational Semantics (ICoS-5)* pages 20–21.
- Niles, Ian and Adam Pease. 2001. Toward a Standard Upper Ontology. In C. Welty and B. Smith, eds., *Proceedings of the 2nd International Conference on Formal Ontology in Information Systems (FOIS-2001)*.
- Niles, Ian and Adam Pease. 2003. Linking lexicons and ontologies: Mapping WordNet to the Suggested Upper Merged Ontology. In *Proceedings of the IEEE International Conference on Information and Knowledge Engineering*, pages 412–416.
- Nivre, Joakim, Marie-Catherine de Marneffe, Filip Ginter, Yoav Goldberg, Jan Hajic, Christopher D. Manning, Ryan McDonald, Slav Petrov, Sampo Pyysalo, Natalia Silveira, Reut Tsarfaty, and Daniel Zeman. 2016. Universal dependencies v1: A multilingual treebank collection. In N. C. C. Chair), K. Choukri, T. Declerck, M. Grobelnik, B. Maegaard, J. Mariani, A. Moreno, J. Odijk, and S. Piperidis, eds., *International Conference on Language Resources and Evaluation (LREC 2016)*. Paris, France: European Language Resources Association (ELRA). ISBN 978-2-9517408-9-1.

- Price, Charlotte, Valeria de Paiva, and Tracy Holloway King. 2008. Context inducing nouns. In *COLING 2008: Proceedings of the workshop on Knowledge and Reasoning for Answering Questions*, pages 9–16. Manchester, UK: Coling 2008 Organizing Committee.
- Rademaker, Alexandre, Valeria de Paiva, Gerard de Melo, Livy Real, and Maira Gatti. 2014. Openwordnet-pt: A project report. In *Proceedings of the 7th Global WordNet Conference (GWC 2014)*.
- Saurí, Roser and James Pustejovsky. 2008. From structure to interpretation: A double-layered annotation for event factuality. In *The 2nd Linguistic Annotation Workshop, LREC, Marrakech. May 26-27, 2008*. Morocco.
- Saurí, Roser and James Pustejovsky. 2009. Factbank: A corpus annotated with event factuality. *Language Resources and Evaluation* 43(3):227–268.
- Saurí, Roser and James Pustejovsky. 2012. Are you sure that this happened? Assessing the factuality degree of events in text. *Computational Linguistics* 38(2).
- Stanovsky, Gabriel, Judith Eckle-Kohler, Yevgeniy Puzikov, Ido Dagan, and Iryna Gurevych. 2017. Integrating deep linguistic features in factuality prediction over unified datasets. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, vol. 2, pages 352–357.
- Tremper, Galina and Anette Frank. 2011. Extending fine-grained semantic relation classification to presupposition relations between verbs. *Bochumer Linguistische Arbeitsberichte* page 129.
- Tremper, Galina and Anette Frank. 2013. A Discriminative Analysis of Fine-Grained Semantic Relations including Presupposition: Annotation and Classification. In S. Dipper, H. Zinsmeister, and B. Webber, eds., *Dialogue and Discourse, Special Issue on Annotating Pragmatic and Discourse Phenomena*, pages 282–322.
- Vossen, Piek, ed. 1998. *EuroWordNet: A Multilingual Database with Lexical Semantic Networks*. Springer. ISBN 0792352955.
- Zaenen, Annie and Lauri Karttunen. 2013. Veridicity annotation in the lexicon? a look at factive adjectives. In *Proceedings of the 9th Joint ISO - ACL SIGSEM Workshop on Interoperable Semantic Annotation*, pages 51–58. Potsdam, Germany: Association for Computational Linguistics.
- Zanzotto, Fabio Massimo and Marco Pennacchiotti. 2010. Expanding textual entailment corpora from wikipedia using co-training. In *Proceedings of the 2nd Workshop on The People’s Web Meets NLP: Collaboratively Constructed Semantic Resources*, pages 28–36. Beijing, China: Coling 2010 Organizing Committee.

